

Pattern Recognition for Neuroimaging Data

Edinburgh, SPM course
April 2013

Overview


- Introduction
 - Univariate & multivariate approaches
 - Data representation
- Pattern Recognition
 - Machine learning
 - Validation & inference
 - Weight maps & feature selection
 - fMRI application
 - Multiclass problem
- Conclusion & PRoNTo

Overview

- Introduction
 - Univariate & multivariate approaches
 - Data representation
- Pattern Recognition
 - Machine learning
 - Validation & inference
 - Weight maps & feature selection
 - fMRI application
 - Multiclass problem
- Conclusion & PRoNTo

Introduction

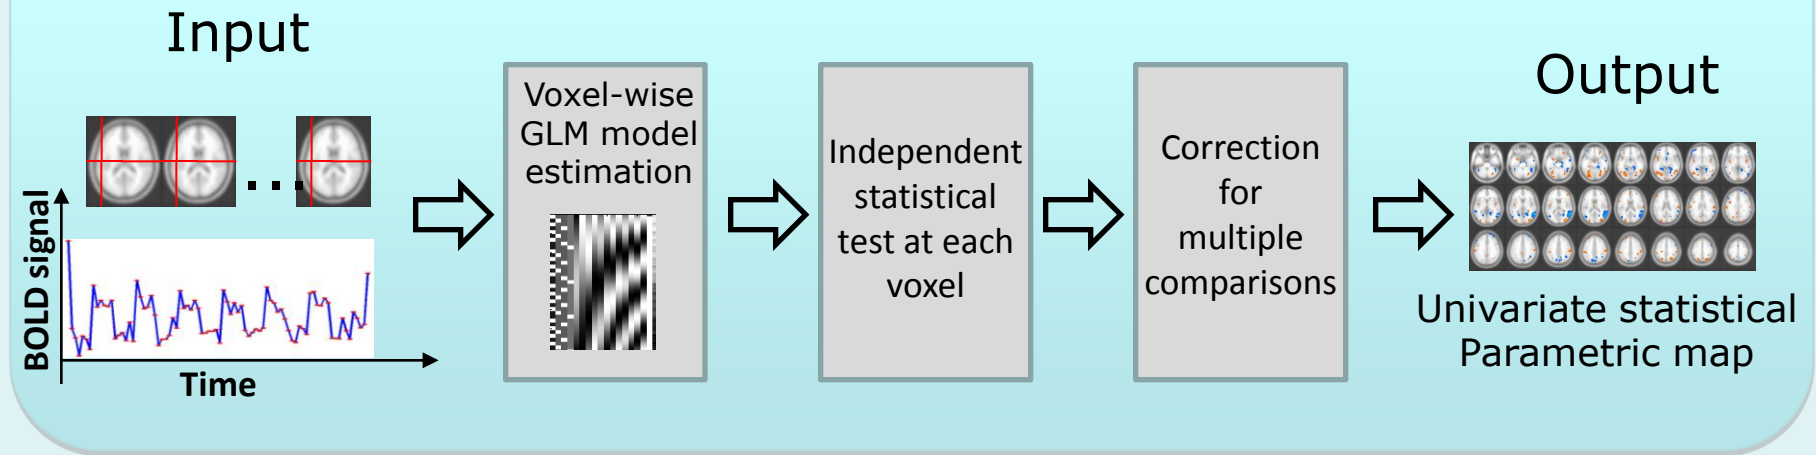
fMRI time series = 4D image
= time series of 3D fMRI's
= 3D array of time series.



Univariate vs. multivariate

Standard univariate approach (SPM)

Standard Statistical Analysis (encoding)

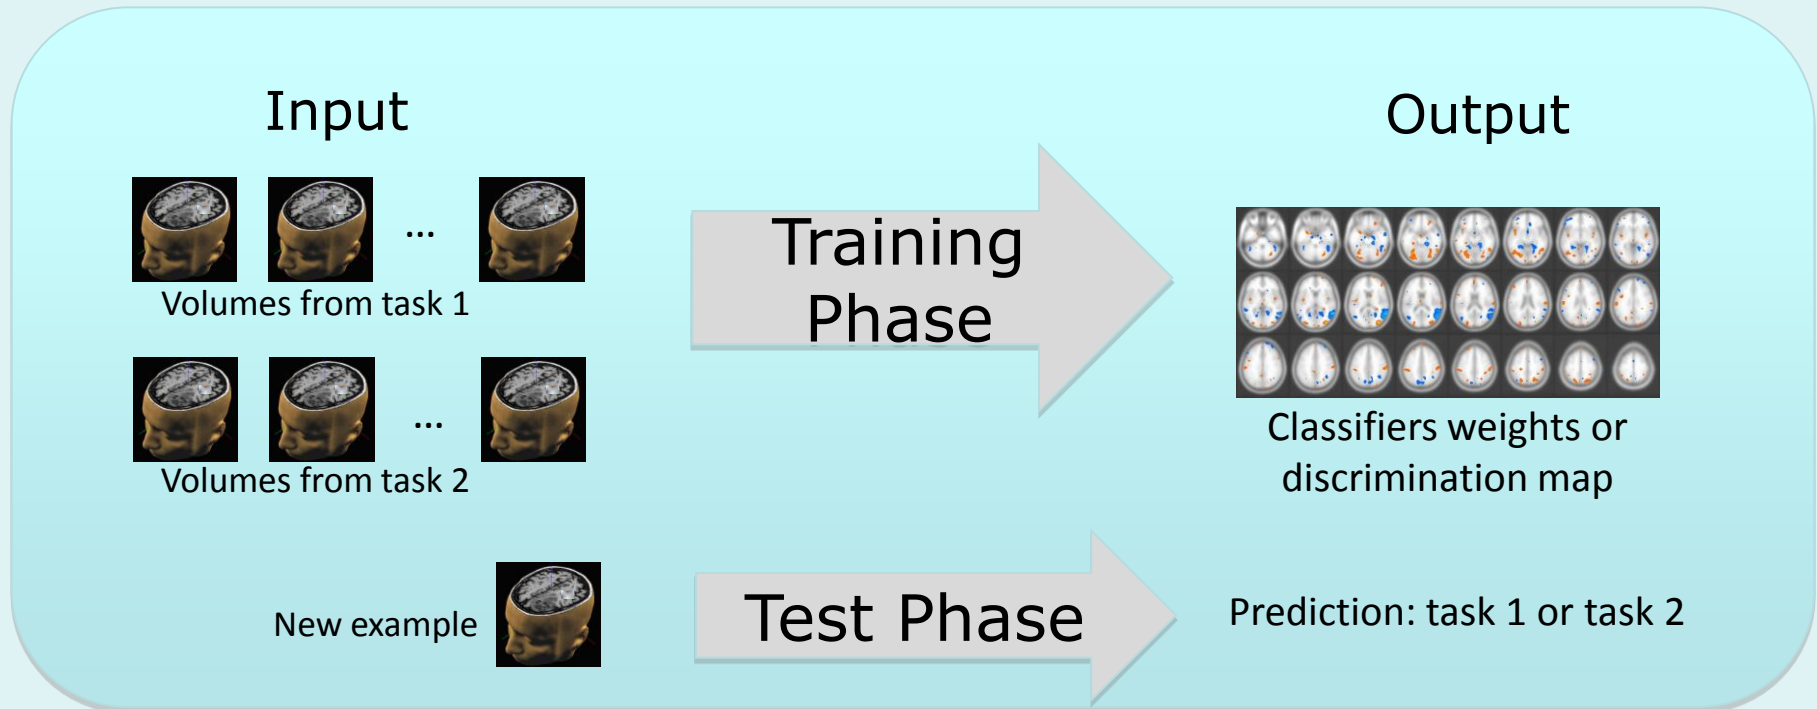


Find the mapping g from explanatory variable X to observed data Y

$$g: X \rightarrow Y$$

Univariate vs. multivariate


Multivariate approach, aka. “pattern recognition”



Find the mapping h from observed data Y to explanatory variable X

$$h: Y \rightarrow X$$

Neuroimaging data



Data dimensions

- dimensionality of a "data point" = #voxels considered
- number of "data point" = #scans/images considered

Note that $\# \text{voxels} \gg \# \text{scans}$!

→ "ill posed problem"


Advantages of pattern recognition

Accounts for the spatial correlation of the data
(multivariate aspect)

- images are multivariate by nature.
- can yield greater sensitivity than conventional (univariate) analysis.

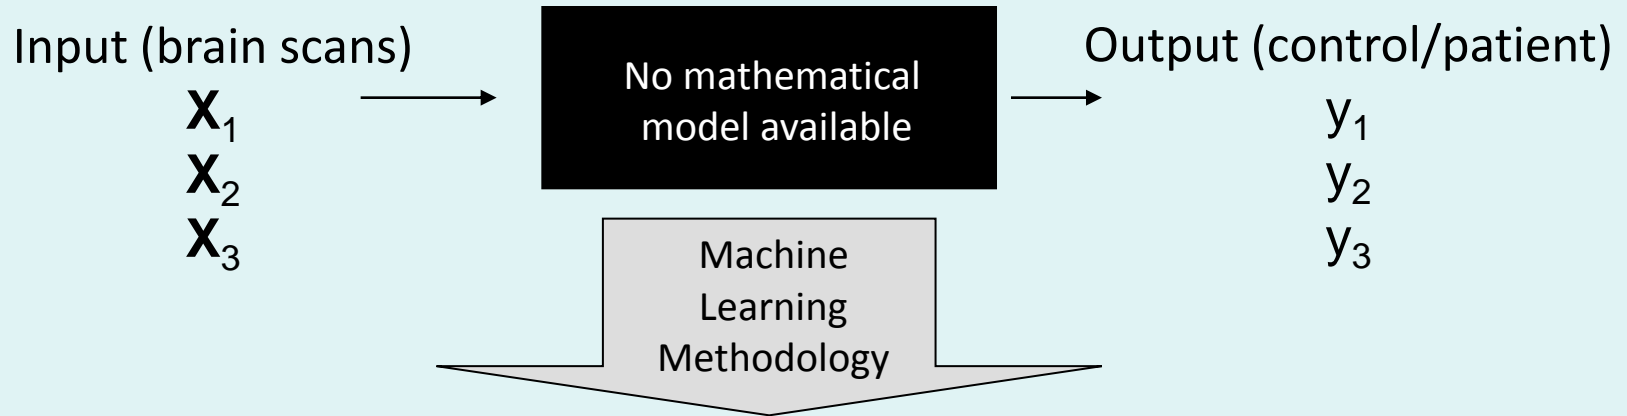
Enable classification/prediction of individual subjects

- 'Mind-reading' or decoding
- Clinical application

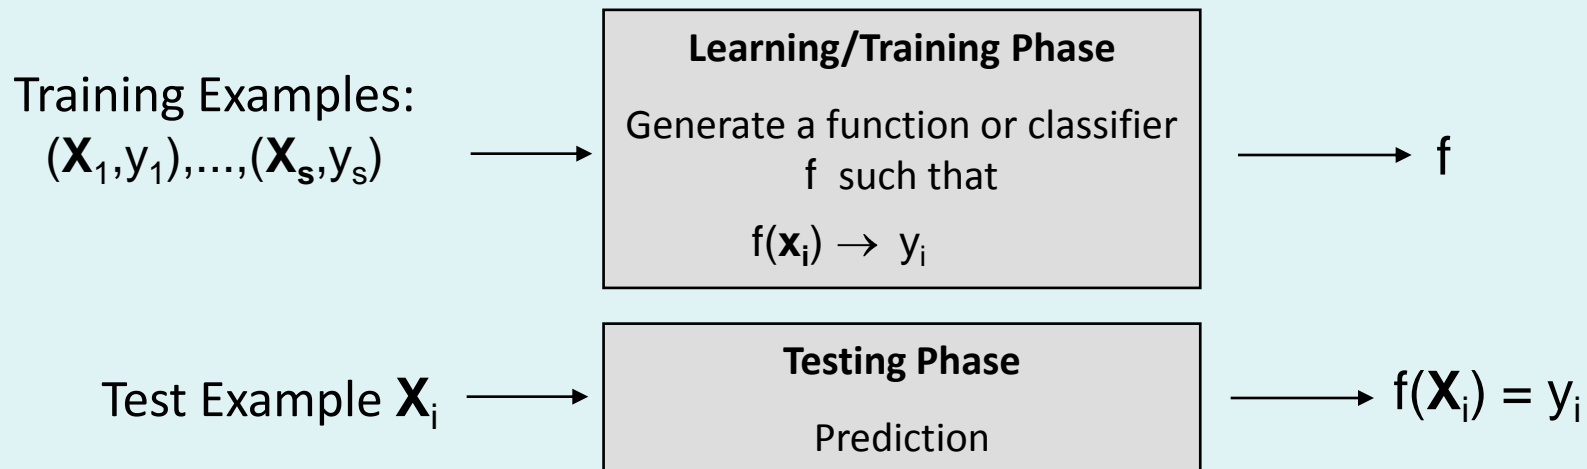


Haynes & Rees, 2006

Pattern recognition framework




Computer-based procedures that learn a function from a *series* of examples



Overview


- Introduction
 - Univariate & multivariate approaches
 - Data representation
- Pattern Recognition
 - Machine learning
 - Validation & inference
 - Weight maps & feature selection
 - fMRI application
 - Multiclass problem
- Conclusion & PRoNTo

Classification example



Note: task1/2 ~ disease/control

Neuroimaging data



Problem: 1000's of features vs. 10's of data points

Possible solutions to dimensionality problem:

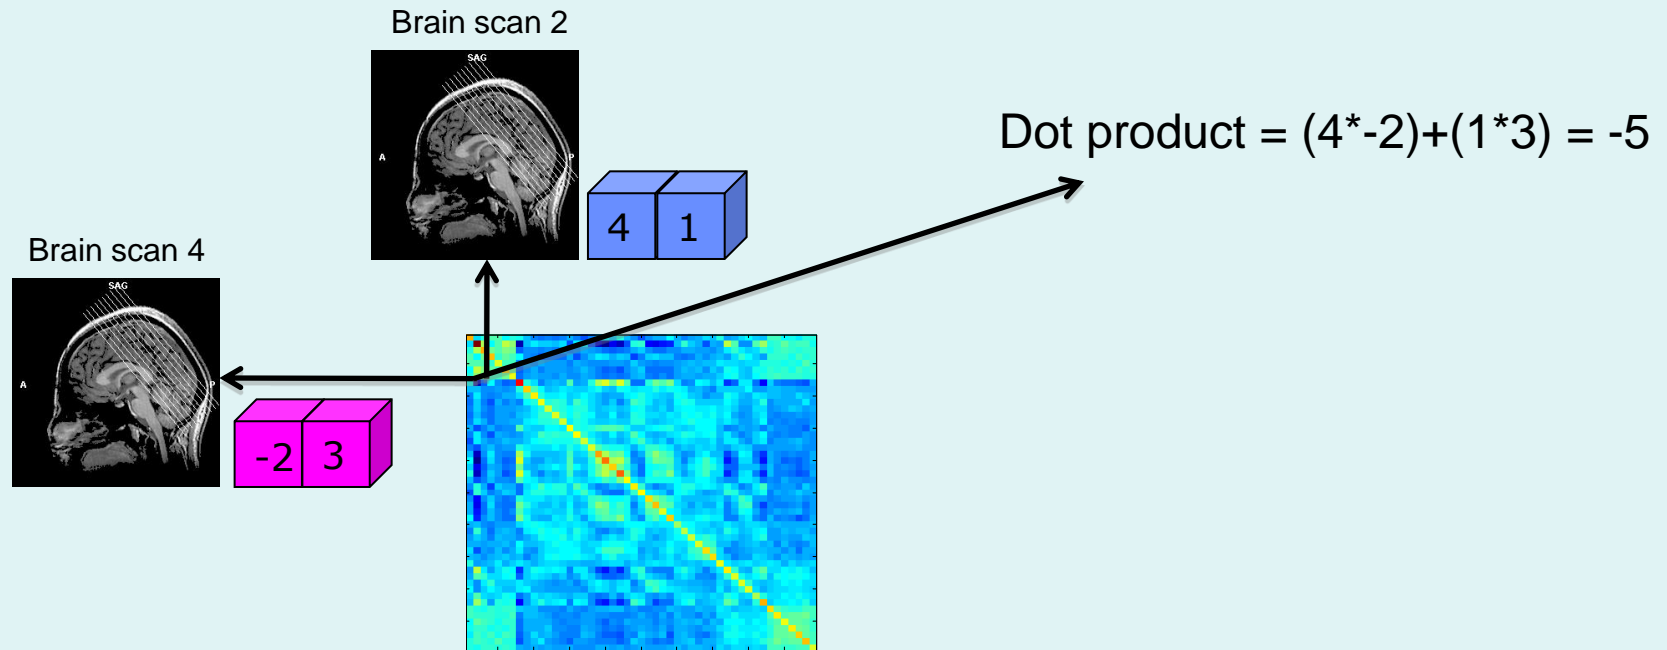
- Feature selection strategies (e.g. ROIS, select only activated voxels)
- (Searchlight)
- **Kernel Methods**

Kernel approaches

- Mathematical trick! → powerful and unified framework (e.g. classification & regression)
- Consist of two parts:
 - build the kernel matrix (mapping into the feature space)
 - train using the kernel matrix (designed to discover linear patterns in the feature space)
- Advantages:
 - computational shortcut → represent linear patterns efficiently in high dimensional space.
 - Using the dual representation with proper regularization → efficient solution of ill-conditioned problems.
- Examples → Support Vector Machine (SVM), Gaussian Processes (GP), Kernel Ridge Regression (KRR),...

Kernel matrix

Kernel matrix = similarity measure



The “kernel function”

- 2 patterns \mathbf{x} and \mathbf{x}^* \rightarrow a real number characterizing their similarity (\sim distance measure).
- simple similarity measure = a dot product \rightarrow linear kernel.

Linear classifier

- hyperplanes through the feature space
- parameterized by
 - a weight vector **w** and
 - a bias term b .
- weight vector **w** = linear combination of training examples \mathbf{x}_i (where $i = 1, \dots, N$ and N is the number of training examples)

$$\mathbf{w} = \sum_{i=1}^N \alpha_i \mathbf{x}_i$$

→ Find the α_i !!!

Linear classifier prediction

General equation: making predictions for a test example \mathbf{x}_* with kernel methods

$$\mathbf{w} = \sum_{i=1}^N a_i \mathbf{x}_i$$

kernel
definition

$$f(\mathbf{x}_*) = \mathbf{w}^\top \mathbf{x}_* + b \longrightarrow \text{Primal representation}$$

$$f(\mathbf{x}_*) = \sum_{i=1}^N a_i \mathbf{x}_i^\top \mathbf{x}_* + b$$


$$f(\mathbf{x}_*) = \sum_{i=1}^N a_i K(\mathbf{x}_i, \mathbf{x}_*) + b \longrightarrow \text{Dual representation}$$

$$f(\mathbf{x}_*) =$$

signed distance to boundary (classification)
predicted score (regression)

Support Vector Machine

SVM = “maximum margin” classifier



$$\mathbf{w} = \sum_{i=1}^N \alpha_i \mathbf{x}_i$$

Support vectors have $\alpha_i \neq 0$


Data: $\langle \mathbf{x}_i, y_i \rangle, i=1, \dots, N$

Observations: $\mathbf{x}_i \in \mathbb{R}^d$


Labels: $y_i \in \{-1, +1\}$

Illustrative example: Classifiers as decision functions


Examples of class 1



Examples of class 2




Weight vector or
Discrimination map



Training

New example



Testing

$$\begin{aligned} f(x) &= (w_1 * v_1 + w_2 * v_2) + b \\ &= (+5 * 0.5 - 3 * 0.8) + 0 \\ &= 0.1 \end{aligned}$$

Positive value
→ Class 1

SVM vs. GP

SVM

- ➔ Hard binary classification
 - simple & efficient, quick calculation but
 - NO 'grading' in output $\{-1, 1\}$

Gaussian Processes


- ➔ probabilistic model
 - more complicated, slower calculation but
 - returns a probability $[0, 1]$
 - can be multiclass

Overview

- Introduction
 - Univariate & multivariate approaches
 - Data representation
- **Pattern Recognition**
 - Machine learning
 - **Validation & inference**
 - Weight maps & feature selection
 - fMRI application
 - Multiclass problem
- Conclusion & PRoNTo

Validation principle

Samples



M-fold cross-validation


- Split data in 2 sets: “train” & “test”
→ evaluation on 1 “fold”



- Rotate partition and repeat
→ evaluations on M “folds”



- Applies to scans/events/blocks/subjects/...
→ Leave-one-out (LOO) approach



Confusion matrix & accuracy

Confusion matrix
= summary table

Accuracy estimation

- Class 0 accuracy, $p_0 = A/(A+B)$
- Class 1 accuracy, $p_1 = D/(C+D)$
- Accuracy, $p = (A+D)/(A+B+C+D)$

Other criteria

- Positive Predictive Value, $PPV = D/(B+D)$
- Negative Predictive Value, $NPV = A/(A+C)$

		$\hat{\omega}$	
		ω_0	ω_1
truth	ω_0	A	B
	ω_1	C	D

Accuracy & Dataset balance

Watch out if #samples/class are different!

Example:

Good overall accuracy (72%) but

- Majority class ($N_1 = 80$), excellent accuracy (90%)
- Minority class ($N_2 = 20$), poor accuracy (0%)

Good practice:

Report

- class accuracies $[p_0, p_1, \dots, p_C]$
- balanced accuracy $p_{\text{bal}} = (p_0 + p_1 + \dots + p_C)/C$

Regression MSE

- LOO error in one fold


$$SE_n = (y_n - f(\mathbf{x}_n))^2$$

- Across all LOO folds

$$R(f, \mathbf{X}) = MSE = \frac{1}{N} \sum_{n=1}^N (y_n - f(\mathbf{x}_n))^2$$

→ Out-of-sample “mean squared error” (MSE)


Other measure:
Correlation between predictions
(across folds!) and ‘true’ targets



Inference by permutation testing

- H_0 : “class labels are non-informative”
- Test statistic = CV accuracy
- Estimate distribution of test statistic under H_0
 - Random permutation of labels
 - Estimate CV accuracy
 - Repeat M times
- Calculate p-value as

$$\frac{1}{M} \sum_m^M (p_m^{perm} \geq p^{real})$$



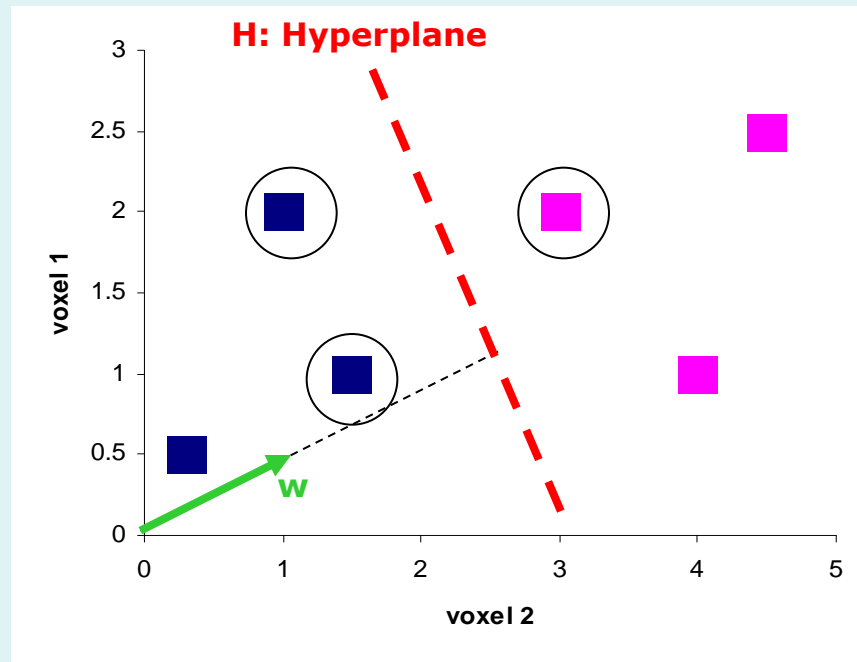




Overview

- Introduction
 - Univariate & multivariate approaches
 - Data representation
- Pattern Recognition
 - Machine learning
 - Validation & inference
 - Weight maps & feature selection
 - fMRI application
 - Multiclass problem
- Conclusion & PRoNTo


Weight vector interpretation

Weight vector

- weight (or discrimination) image !
- how important each voxel is
- for which class "it votes" (mean centred data & $b=0$)

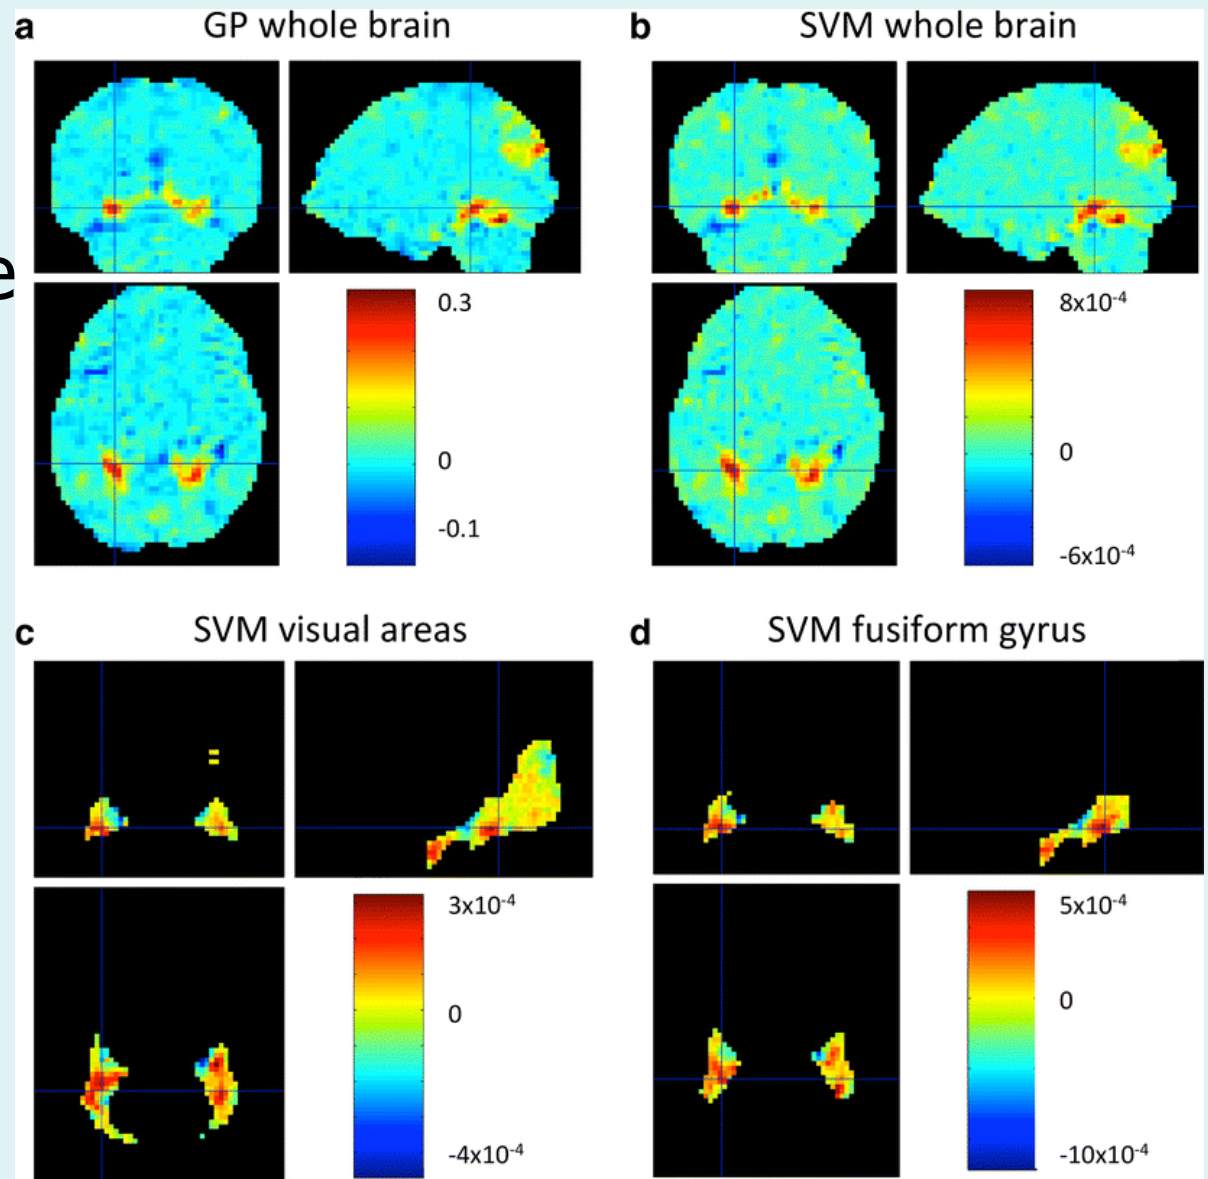


Weight vector
 $W = [0.45 \ 0.89]$
 $b = -2.8$



Example of masks

Linear machine
→ Weight map

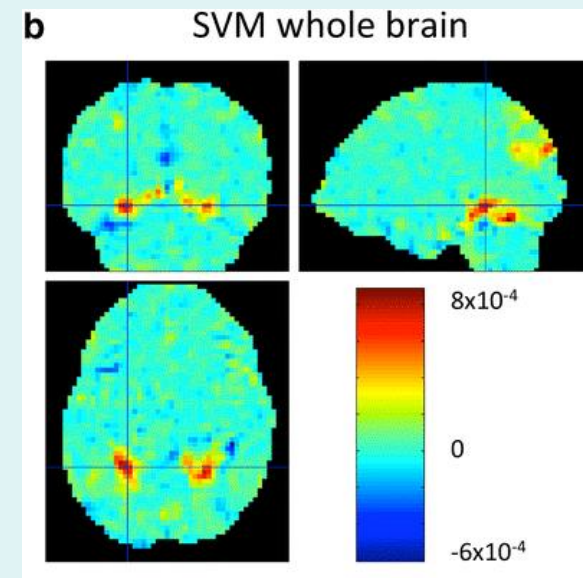


Feature selection

- 1 sample image
 - ➔ 1 predicted value
- use ALL the voxels
 - ➔ NO thresholding of weight allowed!

Feature selection:

- a priori mask
- a priori 'filtering'
- recursive feature elimination/addition
 - ➔ nested cross-validation
(MUST be independent from test data!)



Overview

- Introduction
 - Univariate & multivariate approaches
 - Data representation
- Pattern Recognition
 - Machine learning
 - Validation & inference
 - Weight maps & feature selection
 - **fMRI application**
 - Multiclass problem
- Conclusion & PRoNTo

fMRI designs

Level of inference

- within subject \approx FFX with SPM
→ 'decode' subject's brain state
- between subjects \approx RFX with SPM
→ 'classify' groups, or
regress subjects' parameter

Between subjects

Design

- 2 groups: group A vs. group B
- 1 group: 2 conditions per subject

➔ Extract 1 (or 2) summary image(s) per subject, and classify

Leave-one-out (LOO) cross-validation:

- Leave one subject out (LOSO)
- Leave one subject per group out (LOSGO)


Note: this works for any type of image...

Within subject

Design:

- Block or event-related design
- Accounting for haemodynamic function

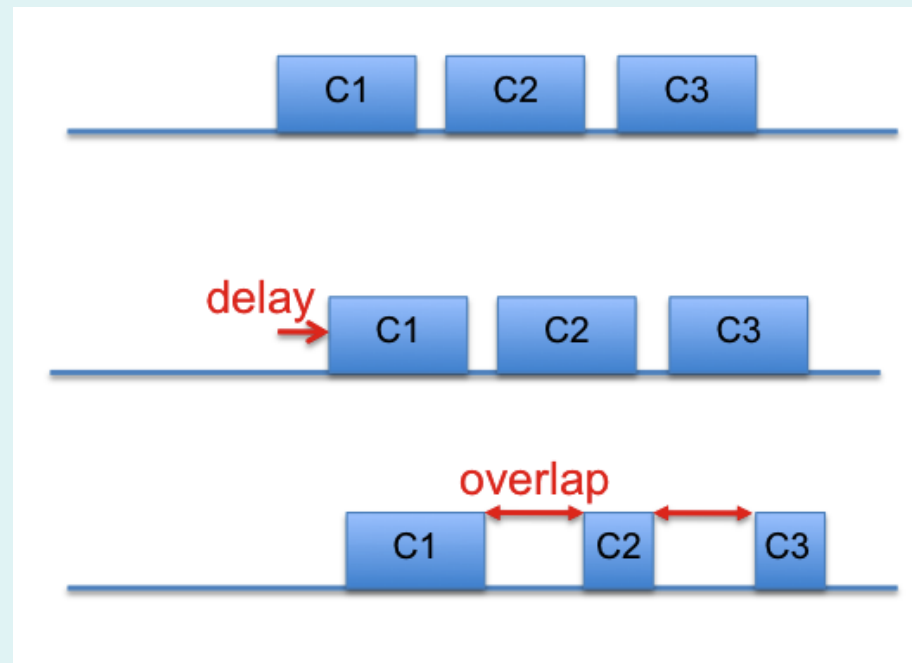


Use single scans



Within subject

Design:

- Block or event-related design
- Accounting for haemodynamic function




Within subject

Design:

- Block or event-related design
- Accounting for haemodynamic function

Averaging/deconvolution



How to?

- Average scans over blocks/events
- Parameter estimate from the GLM with 1 regressor per block/event

Within subject

Design:

- Block or event-related design
- Accounting for haemodynamic function

Leave-one-out (LOO) cross-validation:


- Leave one session/run out
- Leave one block/event out
(danger of dependent data!!!)

Overview


- Introduction
 - Univariate & multivariate approaches
 - Data representation
- Pattern Recognition
 - Machine learning
 - Validation & inference
 - Weight maps & feature selection
 - fMRI application
 - Multiclass problem
- Conclusion & PRoNTo

Multiclass problem


Multiclass machine



Binary machine & one-vs.-others



Binary machine & one-vs.-one



ECOC	SVM codewords			
	C1-C2	C1-C3	C2-C3	L
C1	1	1	0	3
C2	-1	0	1	2
C3	0	-1	-1	1
Example	-1	-1	-1	C3

“Error-Correcting Output Coding” (ECOC) approach

Overview

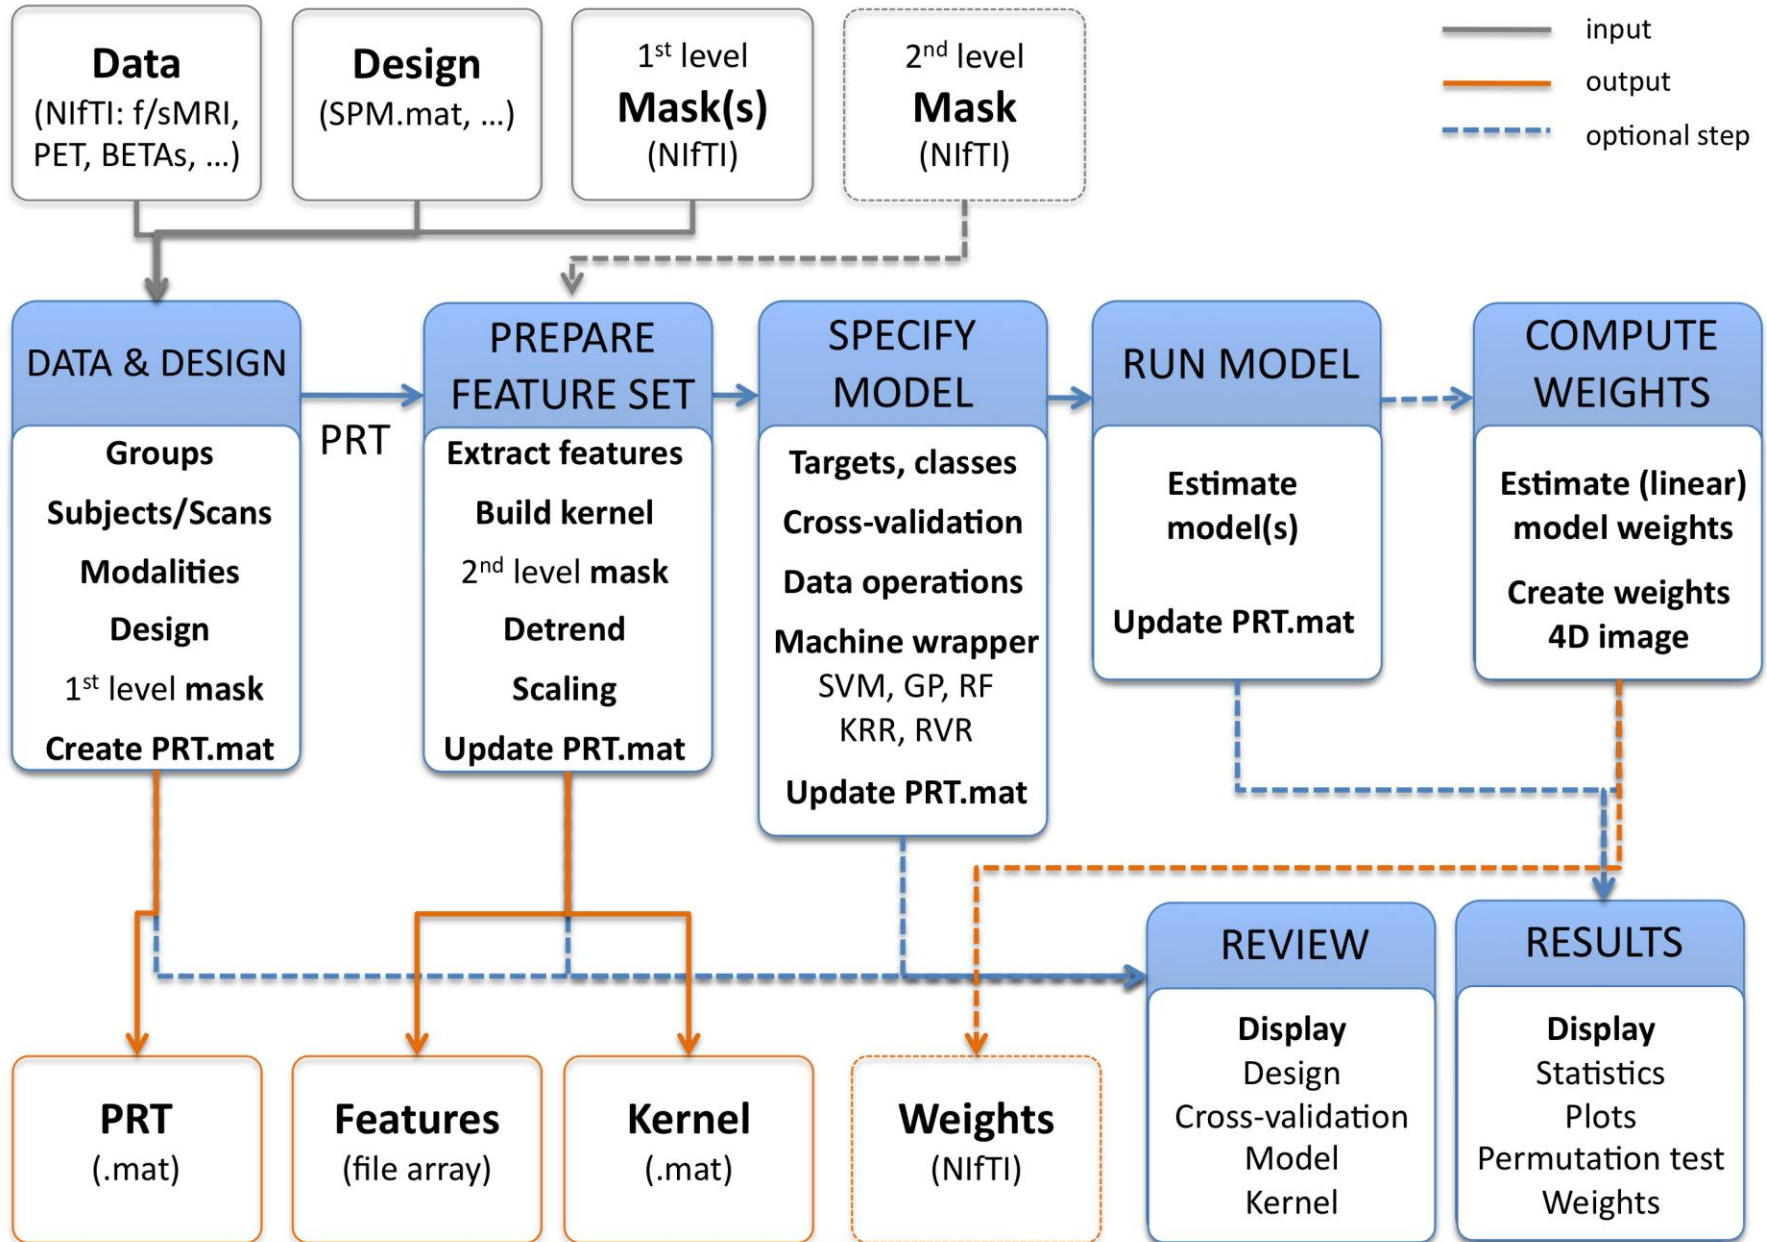
- Introduction
 - Univariate & multivariate approaches
 - Data representation
- Pattern Recognition
 - Machine learning
 - Validation & inference
 - Weight maps & feature selection
 - fMRI application
 - Multiclass problem
- Conclusion & PRoNTo

Conclusions

Key points:

- More sensitivity (~like omnibus test with SPM)
- NO local (voxel/blob) inference
 - ➔ CANNOT report coordinates nor thresholded weight map
- Require cross-validation (split in train/test sets)
 - ➔ report accuracy/PPV (or MSE)
- MUST assess significance of accuracy
 - ➔ permutation approach

PRoNTo FRAMEWORK

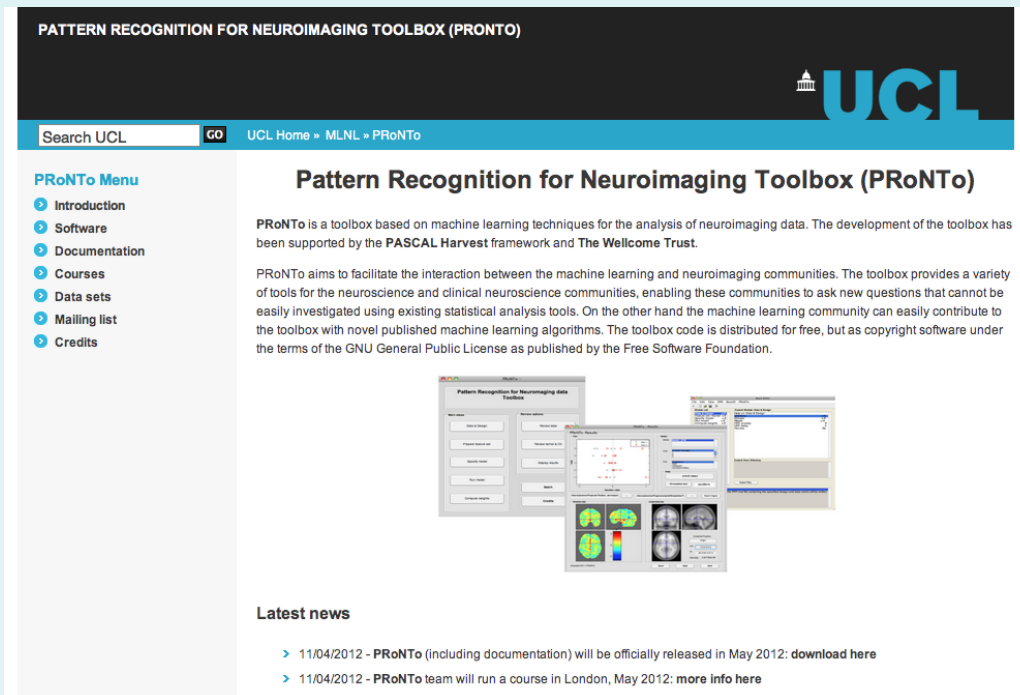


PRoNTTo

“Pattern Recognition for Neuroimaging Toolbox”, aka. PRoNTTo :

<http://www.mlnl.cs.ucl.ac.uk/pronto/>

with references, manual, demo data, course, etc.



Paper: <http://dx.doi.org/10.1007/s12021-013-9178-1>

Thank you for your attention!

Any question?

Thanks to the PRoNTo Team for the borrowed slides. 😊
