

Pattern Recognition for Neuroimaging Data

Edinburgh, SPM course
April 2017



C. Phillips, GIGA – *In Silico* Medicine, ULg, Belgium
c.phillips@ulg.ac.be
<http://www.giga.ulg.ac.be>



Overview


- Introduction
 - Univariate & multivariate approaches
 - Data representation
- Pattern Recognition
 - Machine learning
 - Validation & inference
 - Weight maps & feature selection
 - Applications: groups & fMRI
- Conclusion & Toolboxes

Overview

- Introduction
 - Univariate & multivariate approaches
 - Data representation
- Pattern Recognition
 - Machine learning
 - Validation & inference
 - Weight maps & feature selection
 - Applications: groups & fMRI
- Conclusion & Toolboxes

Introduction

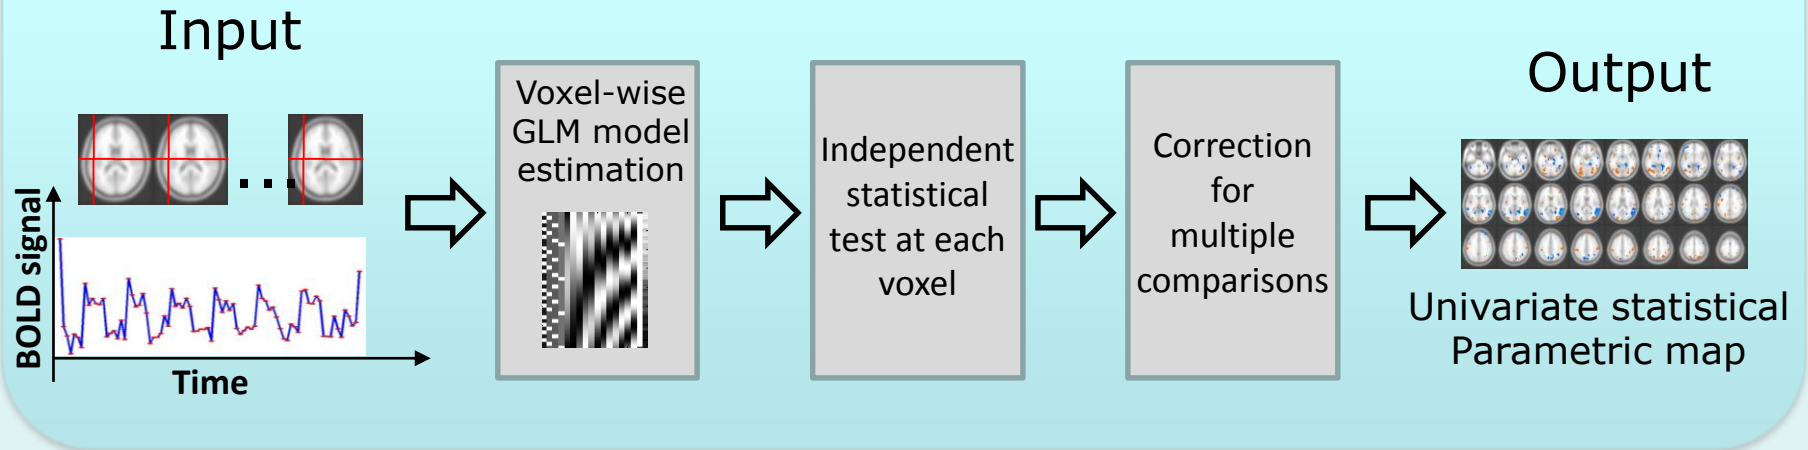
fMRI time series = 4D image
= time series of 3D fMRI's
= 3D array of time series.



Univariate vs. multivariate

Standard univariate approach (SPM)

Standard Statistical Analysis (encoding)

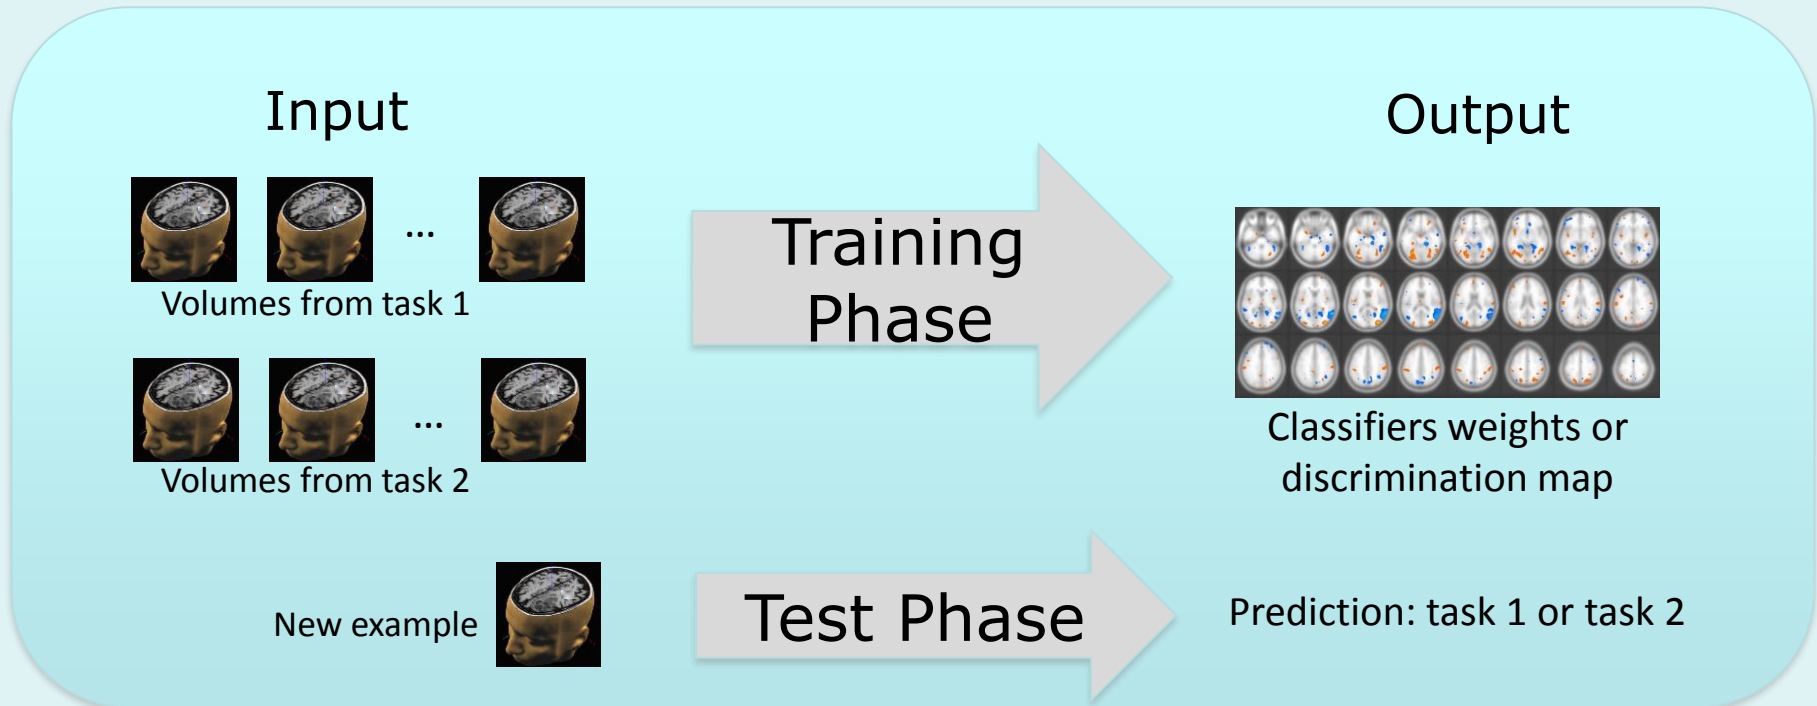


Find the mapping g from explanatory variable X to observed data Y

$$g: X \rightarrow Y$$

Univariate vs. multivariate

Multivariate approach, aka. “pattern recognition”



Find the mapping h from observed data Y to explanatory variable X

$$h: Y \rightarrow X$$

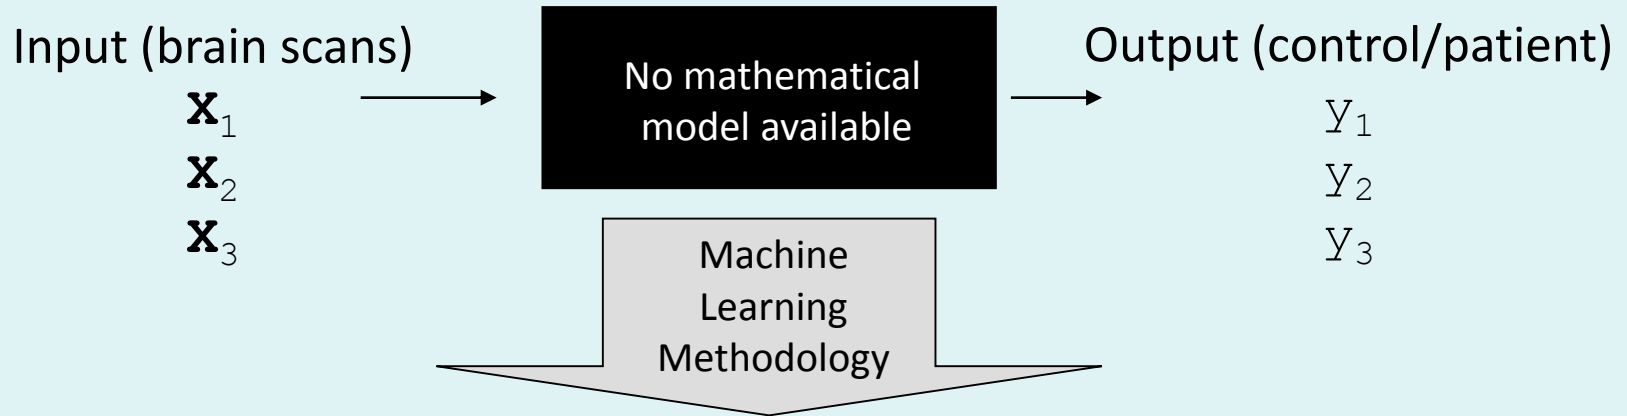
Pattern recognition concept

0	1	2	3	4	5	6	7	8	9	Data X
0	1	2	3	4	5	6	7	8	9	
0	1	2	3	4	5	6	7	8	9	
0	1	2	3	4	5	6	7	8	9	
0	1	2	3	4	5	6	7	8	9	
0	1	2	3	4	5	6	7	8	9	Labels y

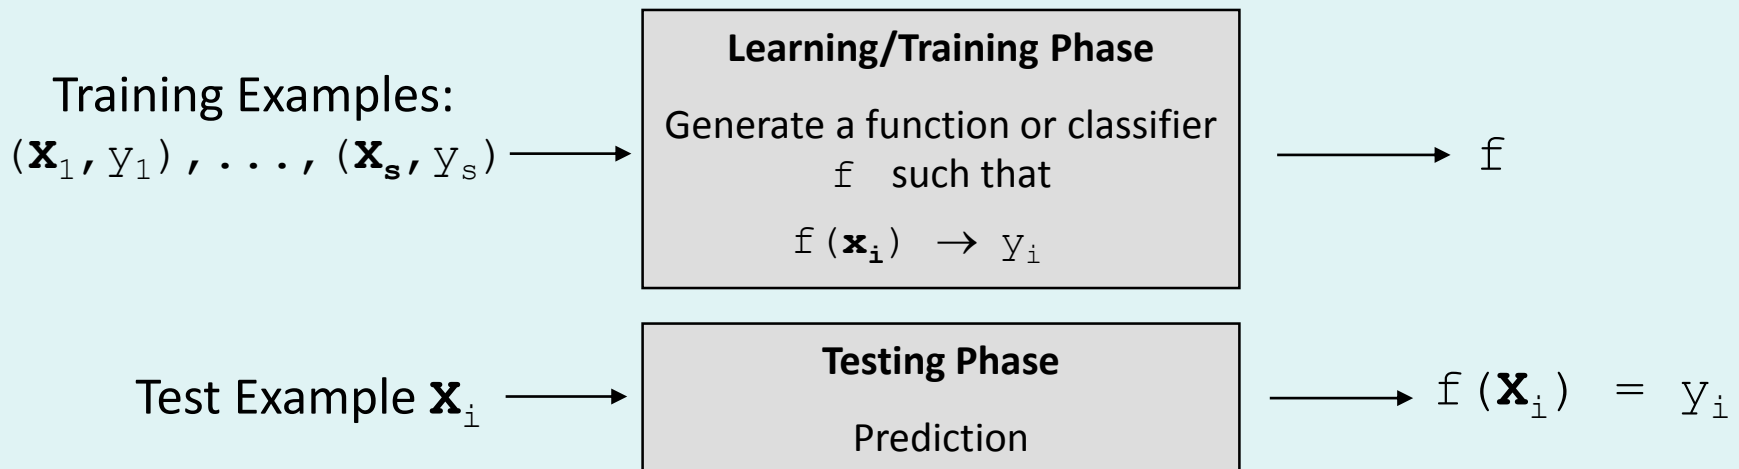
$$f : \mathbf{X} \rightarrow \mathbf{y}$$

$$f : \mathbf{x}_* \rightarrow y_*$$


Pattern recognition framework



Computer-based procedures that learn a function from a *series* of examples



Neuroimaging data



Data dimensions

- dimensionality of a "data point" = #voxels considered
- number of "data point" = #scans/images considered

Note that $\# \text{voxels} \gg \# \text{scans}$!

→ "ill posed problem"

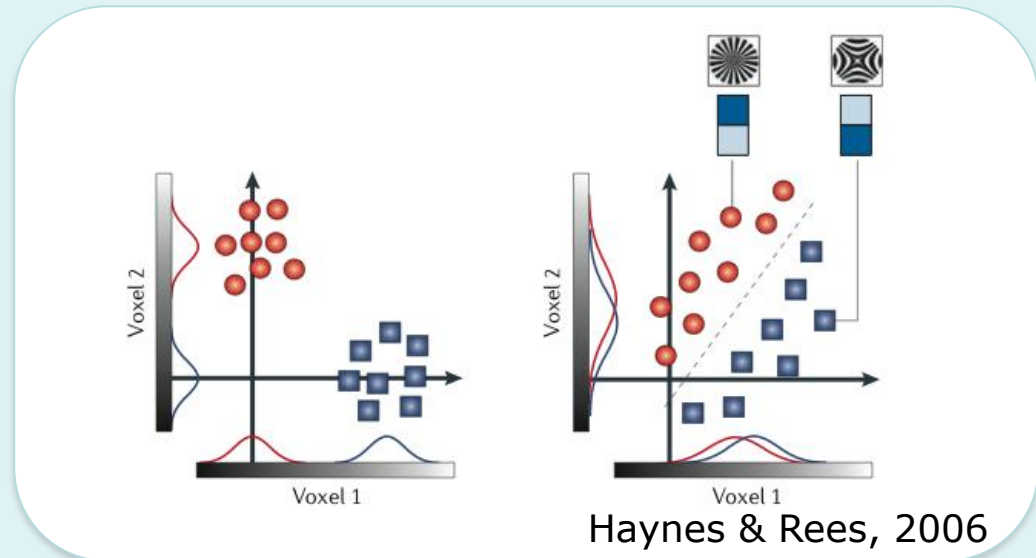
Advantages of pattern recognition

Accounts for the spatial correlation of the data (multivariate aspect)

- images are multivariate by nature.
- can yield greater sensitivity than conventional (univariate) analysis.

Enable classification/prediction of individual subjects

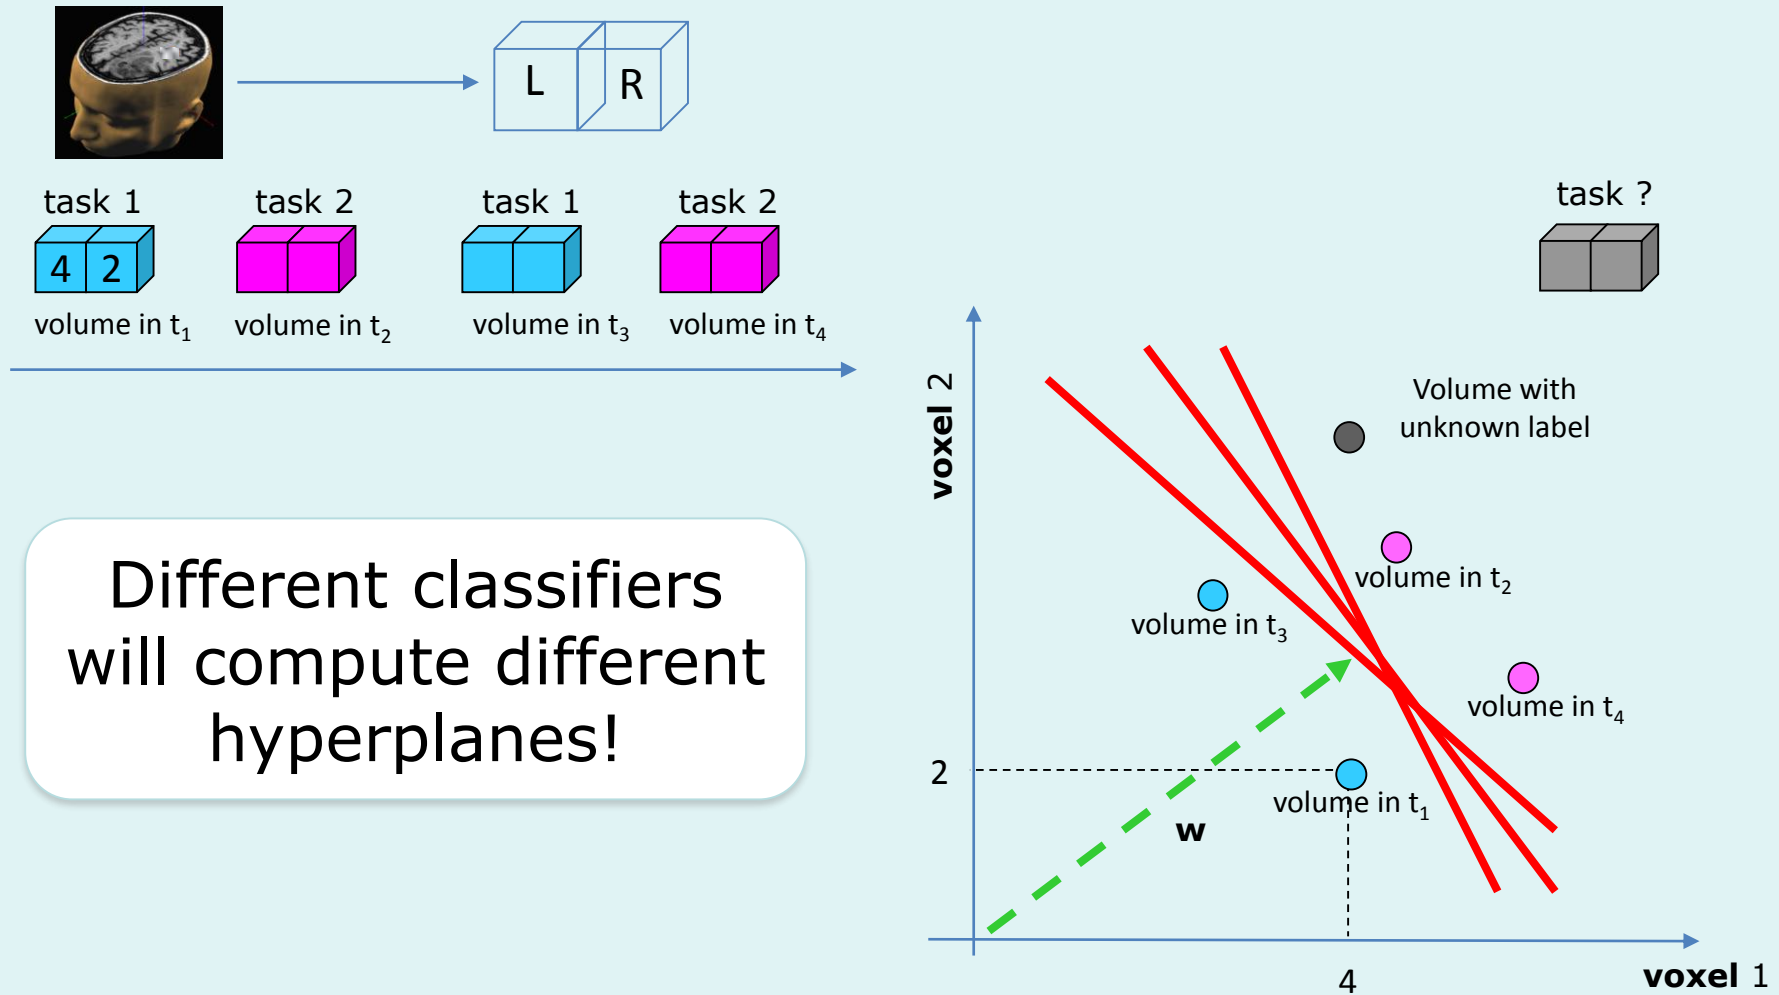
- 'Mind-reading' or decoding applications
- Clinical application



Overview


- Introduction
 - Univariate & multivariate approaches
 - Data representation
- Pattern Recognition
 - Machine learning
 - Validation & inference
 - Weight maps & feature selection
 - Applications: groups & fMRI
- Conclusion & Toolboxes

Classification example



Note: task1/2 ~ disease/control

Neuroimaging data



Problem: 1000's of features vs. 10's of data points

Possible solutions to dimensionality problem:


- Feature selection strategies (e.g. ROIS, select only "activated" voxels)
- (Searchlight)
- **Kernel Methods**

Kernel approaches

- Mathematical trick! → powerful and unified framework (e.g. classification & regression)
- Consist of two parts:
 - build the kernel matrix (mapping into the feature space)
 - train using the kernel matrix (designed to discover linear patterns in the feature space)
- Advantages:
 - computational shortcut → represent linear patterns efficiently in high dimensional space.
 - Using the dual representation with proper regularization → efficient solution of ill-conditioned problems.
- Examples → Support Vector Machine (SVM), Gaussian Processes (GP), Kernel Ridge Regression (KRR),...

Kernel matrix

Kernel matrix = “similarity measure”




The “kernel function”

- 2 patterns \mathbf{x} and \mathbf{x}^* \rightarrow a real number characterizing their similarity (\sim distance measure).
- simple similarity measure = a dot product \rightarrow linear kernel.

Linear classifier

- hyperplane through the feature space
- parameterized by
 - a weight vector \mathbf{w}
(=direction of hyperplane),
 - a bias term b
(=constant term).




- weight vector \mathbf{w} = linear combination of training examples \mathbf{x}_i (where $i = 1, \dots, N$ and N is the number of training examples)

$$\mathbf{w} = \sum_{i=1}^N \alpha_i \mathbf{x}_i$$

→ Find the α_i !!!

Support Vector Machine

SVM = "maximum margin" classifier



$$\mathbf{w} = \sum_{i=1}^N \alpha_i \mathbf{x}_i$$

Support vectors have $\alpha_i \neq 0$


Data: $\langle \mathbf{x}_i, y_i \rangle, i=1, \dots, N$

Observations: $\mathbf{x}_i \in \mathbb{R}^d$

Labels: $y_i \in \{-1, +1\}$

Linear classifier prediction

General equation: making predictions for a test example \mathbf{x}_* with kernel methods




$f(\mathbf{x}_*) =$


signed distance to boundary (classification)
predicted score (regression)

Illustrative example: Classifiers as decision functions


Examples of class 1



Examples of class 2




Weight vector or
Discrimination map



Training

New example



Testing

$$\begin{aligned} f(x) &= (w_1 * v_1 + w_2 * v_2) + b \\ &= (+5 * 0.5 - 3 * 0.8) + 0 \\ &= 0.1 \end{aligned}$$

Positive value
→ Class 1

SVM vs. GP

SVM

- ➔ Hard binary classification
 - simple & efficient, quick calculation but
 - NO 'grading' in output $\{-1, 1\}$

Gaussian Processes


- ➔ probabilistic model
 - more complicated, slower calculation but
 - returns a probability $[0, 1]$
 - can be multiclass

Overview

- Introduction
 - Univariate & multivariate approaches
 - Data representation
- Pattern Recognition
 - Machine learning
 - **Validation & inference**
 - Weight maps & feature selection
 - Applications: groups & fMRI
- Conclusion & Toolboxes

Validation principle

Samples = {variables, labels}



M-fold cross-validation


- Split data in 2 sets: “train” & “test”
→ evaluation on 1 “fold”



- Rotate partition and repeat
→ evaluations on M “folds”



- Applies to scans/events/blocks/subjects/...
→ Leave-one-out (LOO) approach



Confusion matrix & accuracy

Confusion matrix
= summary table

Accuracy estimation

- Class 0 accuracy, $p_0 = A/(A+B)$
- Class 1 accuracy, $p_1 = D/(C+D)$
- Accuracy, $p = (A+D)/(A+B+C+D)$

Other criteria

- Positive Predictive Value, $PPV = D/(B+D)$
- Negative Predictive Value, $NPV = A/(A+C)$

		$\hat{\omega}$	
		ω_0	ω_1
truth	ω_0	A	B
	ω_1	C	D

Accuracy & Dataset balance

Watch out if #samples/class are different!!!

Example: Classes A/B with 80/20 samples each

- observed $\mathbf{a_{tot} = 70\%}$ overall accuracy but
 - within class A ($N_A = 80$), excellent accuracy (85%)
 - within class B ($N_B = 20$), poor accuracy (10%)
- balanced accuracy $\mathbf{a_{bal} = 47,5\%}$!

Good practice:

Report

- class accuracies $[a_0, a_1, \dots, a_C]$
- balanced accuracy $a_{bal} = (a_0 + a_1 + \dots + a_C) / \#Classes$

Regression validation

“Mean squared error” (MSE):

- LOO error in one fold

$$SE_n = (y_n - f(\mathbf{x}_n))^2$$

- Across all LOO folds


$$R(f, \mathbf{X}) = MSE = \frac{1}{N} \sum_{n=1}^N (y_n - f(\mathbf{x}_n))^2$$

→ Out-of-sample “mean squared error” (MSE)

Other measure:

Correlation between:

- predictions (across folds!), and
- ‘true’ targets




Inference by permutation testing

- H_0 : “class labels are non-informative”
- Test statistic = CV accuracy (total or balanced)
- Estimate distribution of test statistic under H_0
 - ➔ Random permutation of labels
 - ➔ Estimate accuracy
 - ➔ Repeat M times
- Calculate p-value

as

$$p = \frac{1}{M} \sum_m^M (a_m^{\text{perm}} \geq a^{\text{true}})$$









Overview

- Introduction
 - Univariate & multivariate approaches
 - Data representation
- Pattern Recognition
 - Machine learning
 - Validation & inference
 - Weight maps & feature selection
 - Applications: groups & fMRI
- Conclusion & Toolboxes


Weight vector interpretation

Weight vector

- weight (or discrimination) image !
- how important each voxel is
- for which class "it votes" (mean centred data & $b=0$)



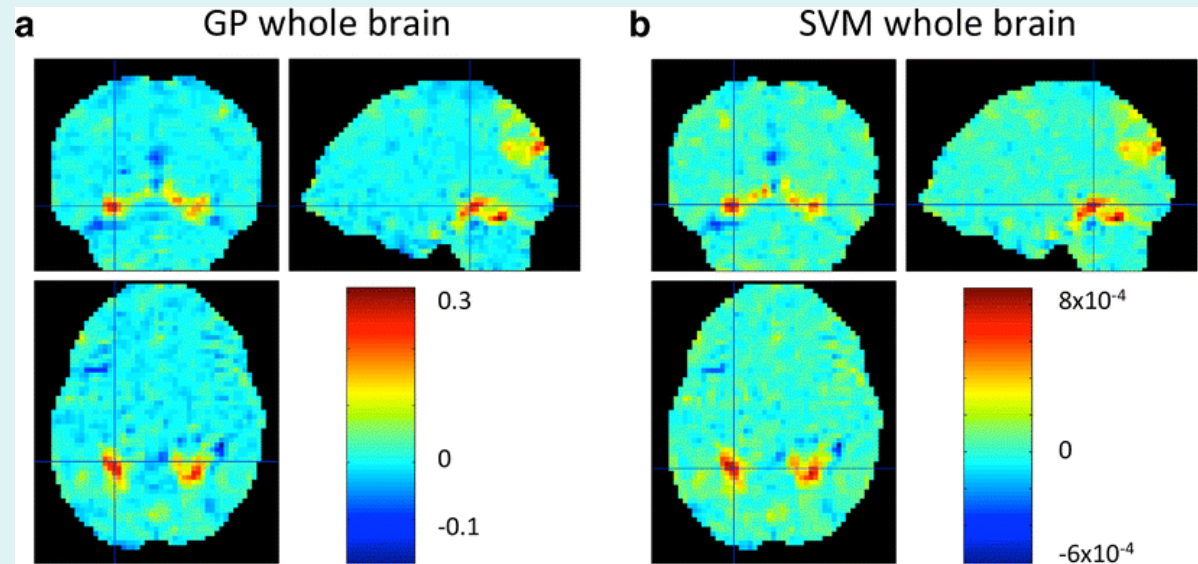
Weight vector
 $W = [0.89 \ 0.45]$
 $b = -2.8$



Weight maps for different masks

Linear machine
→ Weight map

Different mask
→ different
feature set
→ different
weight map




Feature selection

- 1 sample image
 - 1 predicted value
- use ALL the voxels
 - NO thresholding of weight allowed!

Feature selection:

- *a priori* mask
 - *a priori* 'filtering'
 - Multiple Kernel Learning
 - (Search Light)
 - Recursive Feature Elimination/Addition
- MUST be independent from test data!



Overview

- Introduction
 - Univariate & multivariate approaches
 - Data representation
- Pattern Recognition
 - Machine learning
 - Validation & inference
 - Weight maps & feature selection
 - Applications: groups & fMRI
- Conclusion & Toolboxes

Application & designs

Levels of “inference”

- within subject \approx FFX with SPM
 - ➔ ‘decode’ subject’s brain states
 - ➔ multiple images: typically fMRI
- between subjects \approx RFX with SPM
 - ➔ ‘classify’ groups, e.g. patients vs. controls
(or regress subjects’ parameter)
 - ➔ 1 (or few) image(s)/subject : contrast (fMRI),
structural (GM, FA, Jacobian,...), PET,
etc.

Between subjects

Design

- 2 groups: group A vs. group B
- 1 group: 2 conditions per subject (e.g. before/after treatment)
- 1 group: 1 target score

➔ Extract 1 (or a few) summary image(s) per subject, and classify/regress


Leave-one-out (LOO) cross-validation:

- Leave one subject out (LOSO)
- Leave one subject per group out (LOSGO)

Within subject, fMRI

Design:


- Block or event-related design?
- How to account for haemodynamic function?



Within subject, fMRI

Design:

- Block design
- Use single scans






C1: Condition 1
C2: Condition 2
BL: Baseline

Within subject, fMRI

Design:


- Block or event-related design
- Accounting for haemodynamic function, with HRF 'delay' & 'overlap'




Within subject, fMRI

Design:

- Block or event-related design
- Accounting for haemodynamic function by averaging/deconvolution



C1: Condition 1
C2: Condition 2
BL: Baseline



How to?

- **Average scans** over blocks/events
- **Parameter (beta) estimate** from the GLM with 1 regressor per block/event

Within subject

Design:

- Block or event-related design
- Accounting for haemodynamic function

Leave-one-out (LOO) cross-validation:

- Leave one session/run out
- Leave one block/event out
(danger of dependent data!!!)

Overview

- Introduction
 - Univariate & multivariate approaches
 - Data representation
- Pattern Recognition
 - Machine learning
 - Validation & inference
 - Weight maps & feature selection
 - Applications: groups & fMRI
- Conclusion & Toolboxes

Conclusions

Key points:

- NO local (voxel/blob) inference
 - ➔ CANNOT report coordinates nor thresholded weight map
- Require cross-validation (split in train/test sets)
 - ➔ report accuracy/PPV (or MSE)
- MUST assess significance of accuracy
 - ➔ permutation approach
- Could expect more sensitivity (~like omnibus test with SPM)
- Different questions & Different designs!?

“Univariate vs. multivariate” concepts

Univariate

- 1 voxel
- target \rightarrow data
- look for difference or correlation
- General Linear Model
- GLM inversion
- calculate contrast of interest

Multivariate

- 1 volume
- data \rightarrow target
- look for similarity or score
- Specific machine (SVM, GP,...)
- training & testing cross-validation
- estimate accuracy of prediction

Existing toolboxes

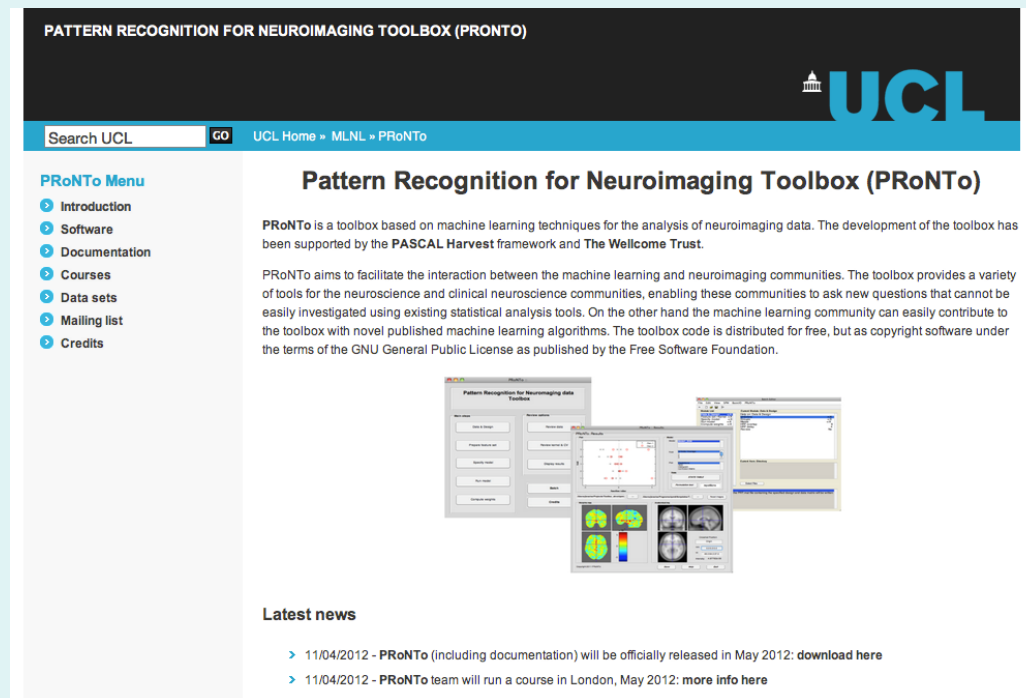
- ✓ Princeton Multi-Voxel Pattern Analysis (MVPA) Toolbox (<http://code.google.com/p/princeton-mvpa-toolbox/>)
 - (+) Matlab based, AFNI/BV oriented but some SPM5 compatibility,
 - (-) only(?) neural network classification, only(?) fMRI data, last version from 2009.
- ✓ Multivariate Pattern Analysis with Python – PyMVPA (<http://www.pymvpa.org/>)
 - (+) looks powerful with loads of tools,
 - (-) Python based, OS specific (Linux).
- ✓ The Decoding Toolbox – TDT (<https://sites.google.com/site/tdtdecodingtoolbox/>)
 - (+) search-light & SVM, ROI/whole brain, feature selection, flexible CV & design, SPM8/12 compatibility
 - (-) no SPM-style batching (?)
- ✓ Brain Voyager MVPA – BV-MVPA (<http://www.brainvoyager.com/bvqx/>)
 - (+) search-light & SVM,
 - (-) proprietary code, fMRI only, no way of adding own machine

PRoNTo

Pattern Recognition for Neuroimaging Toolbox

<http://www.mlnl.cs.ucl.ac.uk/pronto/>

with references, manual, demo data, course, etc.



Schrouff et al, 2013.

Afternoon workshop

More about

- Weight interpretation
- Machines & “multi-kernel learning”
- Nested CV & parameter optimization
- Feature selection & search light approach
- Multiclass ($N > 2$) classification
- fMRI design & group analysis
- ...

And practical demo of PRoNTo:

- fMRI & group analysis
- GUI and batching

Thank you for your attention!

Any question?

Thanks to the PRoNTo Team for the borrowed slides. 😊

References

- C. Bishop. Pattern Recognition and Machine Learning. Springer, 2006.
- Haynes and Rees (2006). Decoding mental states from brain activity in humans. *Nat. Rev. Neurosci.* 7, 523–534.
- Mourão-Miranda et al. (2006). The impact of temporal compression and space selection on SVM analysis of single-subject and multi-subject fMRI data. *Neuroimage* 33, 1055–1065.
- Noirhomme, et al. (2014). Biased binomial assessment of cross-validated estimation of classification accuracies illustrated in diagnosis predictions. *Neuroimage Clin.* 4, 687–694.
- Pereira et al., Machine learning classifiers and fMRI: A tutorial overview, *NeuroImage* 45, 2009.
- C. Rasmussen and C. K. I. Williams. Gaussian Processes for Machine Learning . The MIT Press, 2006.
- Schrouff et al. (2013). PRoNTo: Pattern Recognition for Neuroimaging Toolbox. *Neuroinformatics*, 11, 319–337.
- J. Shawe-Taylor and N. Christianini (2004) Kernel Methods for Pattern Analysis. Cambridge: Cambridge University Press.
